Computer Network Design

Network Layer II

Yalda Edalat – Spring 23

Review

- Which devices are on same networks?
- A: 192.168.137.44/21
- B: 192.168.141.12/21
- C: 192.168.145.3/21

- A: 11000000.10101000.10001001.00101100
- B: 11000000.10101000.10001101.00001100
- C: 11000000.10101000.10010001.00000011

IP Forwarding

- IP addresses on one network belong to the same prefix
- Node uses a table that lists the next hop for IP prefixes

Prefix	Next Hop
192.24.0.0/18	D
192.24.12.0/22	B

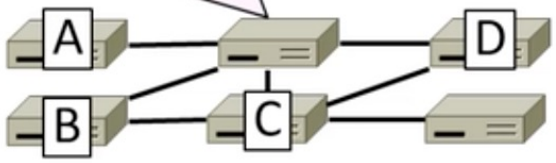

Longest Matching Prefix

- Prefixes in the table might overlap!

- Longest matching prefix forwarding rule:
  - For each packet, find the longest prefix that contains the destination address, i.e., the most specific entry
  - Forward the packet to the next hop router for that prefix

Longest Matching Prefix (2)

Prefix	Next Hop
192.24.0.0/18	D
192.24.12.0/22	B

192.24.6.0 →
192.24.14.32 →
192.24.54.0 →

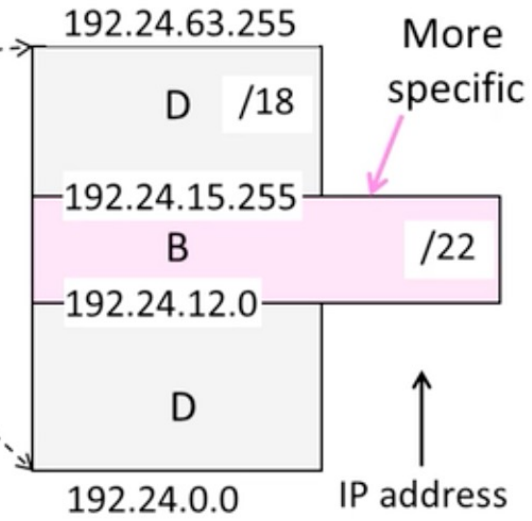

Host/Router Distinction

- In the Internet:
  - Routers do the routing, know which way to all destinations
  - Hosts send remote traffic (out of prefix) to nearest router

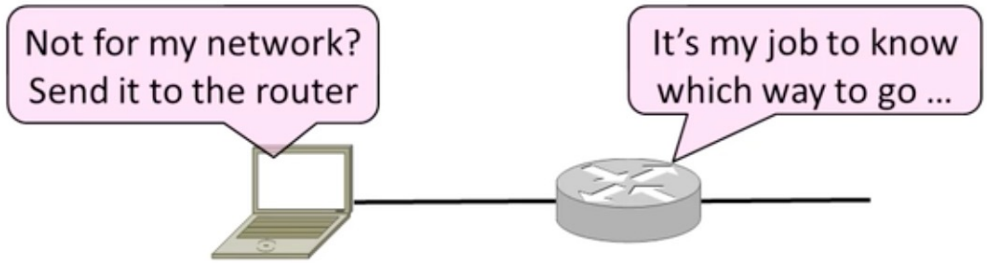

Host Forwarding Table

- Give using longest matching prefix
  - 0.0.0.0/0 is a default route that catches all IP addresses

Prefix	Next Hop
My network prefix	Send direct to that IP
0.0.0.0/0	Send to my router

Example

- Source: 192.168.22.78/24
- Destination: 192.168.23.71/24

- What is next hop? Is it 192.168.23.71? Or default router of source?

Getting IP Addresses

- Problem:
  - A node wakes up for the first time …
  - What is its IP address? What is the IP address of its router? Etc.
  - At least Ethernet address is on NIC

Getting IP Addresses (2)

1. Manual configuration (old days)
   - Can't be factory set, depends on use

2. A protocol for automatically configuring addresses (DHCP)
   - Shifts burden from users to IT folk

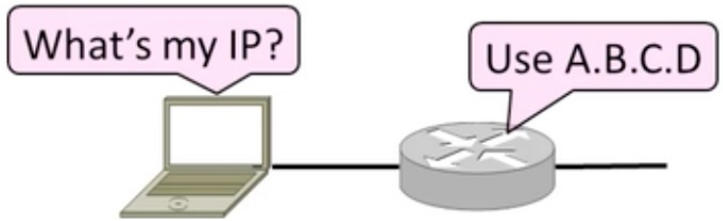

DHCP

- DHCP (Dynamic Host Configuration protocol) is widely used

- It leases IP address to nodes
- Provide other parameters too
  - Network prefix
  - Address of local router (default router)
  - DNS server, time server, etc.

- DHCP is a client-server application
  - Uses UDP ports 67, 68

DHCP Addressing

- Bootstrap issue:
  - How does node send a message to DHCP server before it is configured?

- Answer:
  - Node sends broadcast messages that delivered to all nodes on the network
  - Broadcast address is all 1s (IP: 255.255.255.255) (Ethernet: ff:ff:ff:ff:ff:ff)

DHCP Messages

- To renew an existing lease, an abbreviated sequence is used:
  - REQUEST, followed by ACK

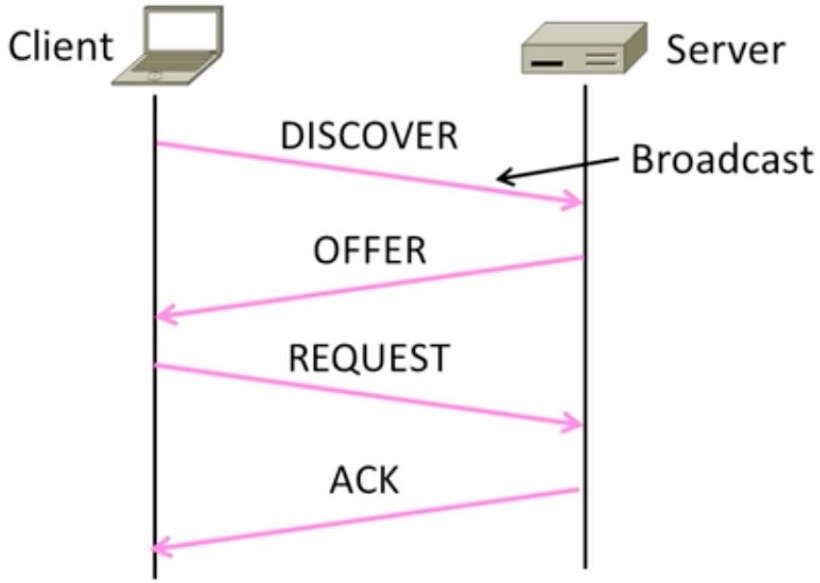

Sending an IP Packet

- Problem:
  - A node needs link layer addresses to send a frame over the local link
  - How does it get the destination link address from a destination IP address?

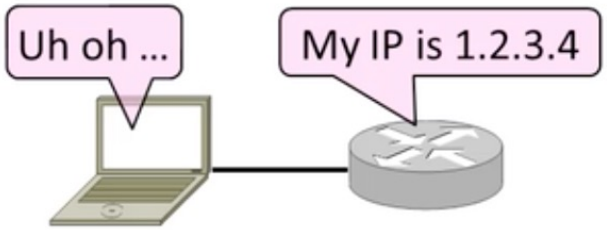

ARP (Address Resolution Protocol)

- Node uses to map a local IP address to its link layer addresses

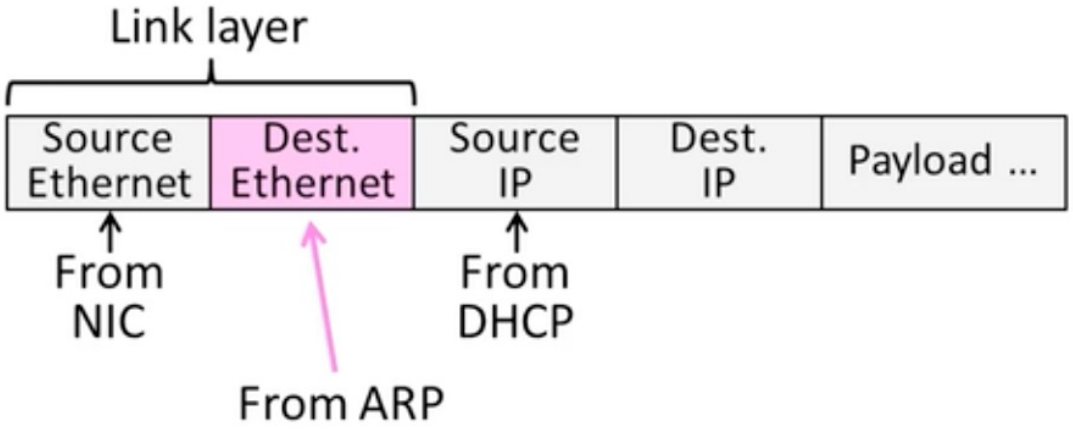

ARP Protocol Stack

- ARP sits right on top of link layer
  - No servers, just asks node with target IP to identify itself
  - Uses broadcast to reach all nodes

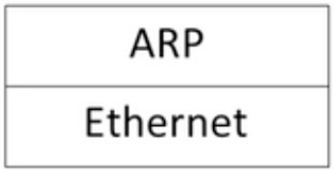

ARP Messages

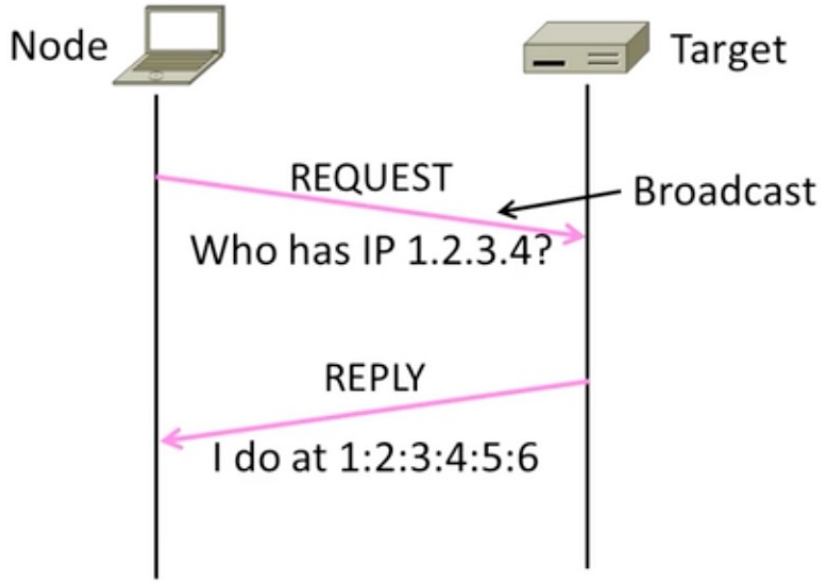

Packet Size Problem

- Different networks have different maximum packet sizes (MTU)
- How do we connect networks with different maximum packet sizes?
  - Need to split up packets or discover the largest size to use

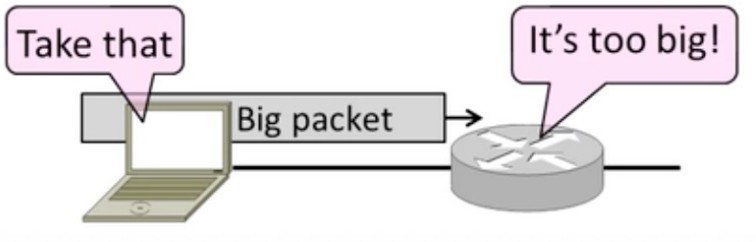

- Prefer large packets for efficiency
  - But what size is too large?
  - Difficult because node doesn't know complete network path

Packet Size Solution

- Fragmentation
  - Split up large packets in the network if they are too large to send
  - Classic method, dated
- Discovery
  - Find the largest packet that fits on the network path and use it
  - IP uses today instead of fragmentation

IPv4 Fragmentation

- Routers fragment packets that are too large to forward
- Receiving host reassembles to reduce load on routers

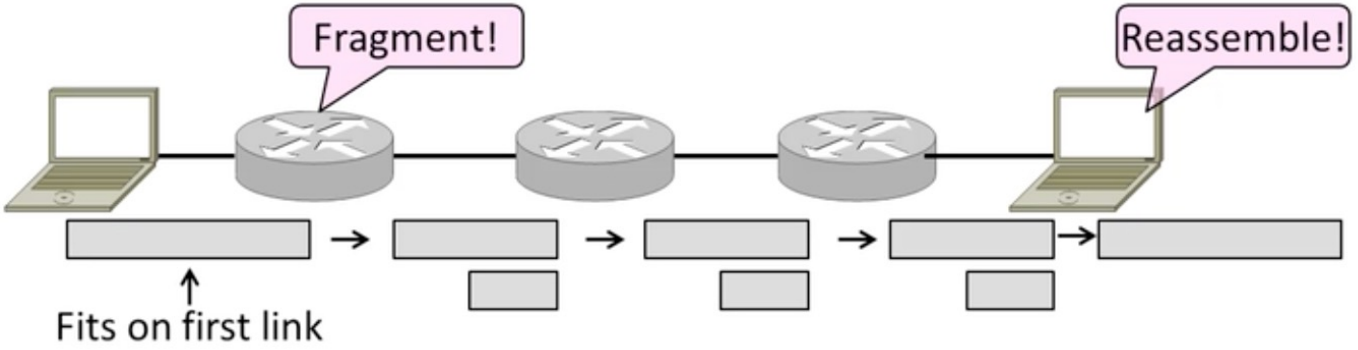

IPv4 Fragmentation Fields

- Header fields used to handle packet size differences
  - Identification, Fragment offset, MF/DF control bits

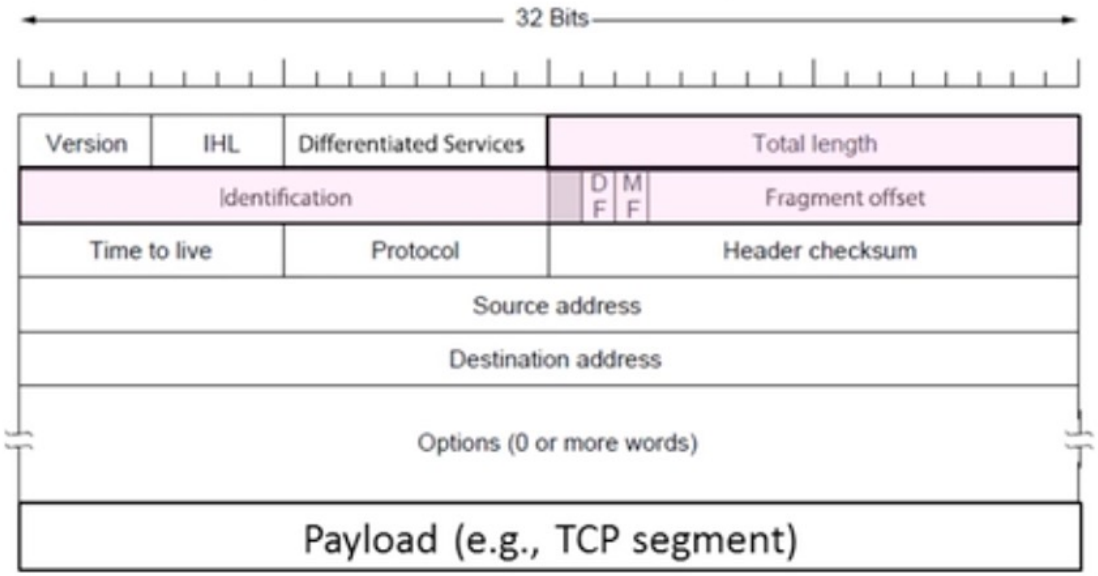

IPv4 Fragmentation Procedure

- Routers split a packet that is too large:
  - Typically break into large pieces
  - Copy IP header to pieces
  - Adjust length on pieces
  - Set offset to indicate position
  - Set MF (More Fragments on all pieces except last

- Receiving hosts reassembles the pieces:
  - Identification field links pieces together, MF tells if it has all pieces

IPv4 Fragmentation (2)

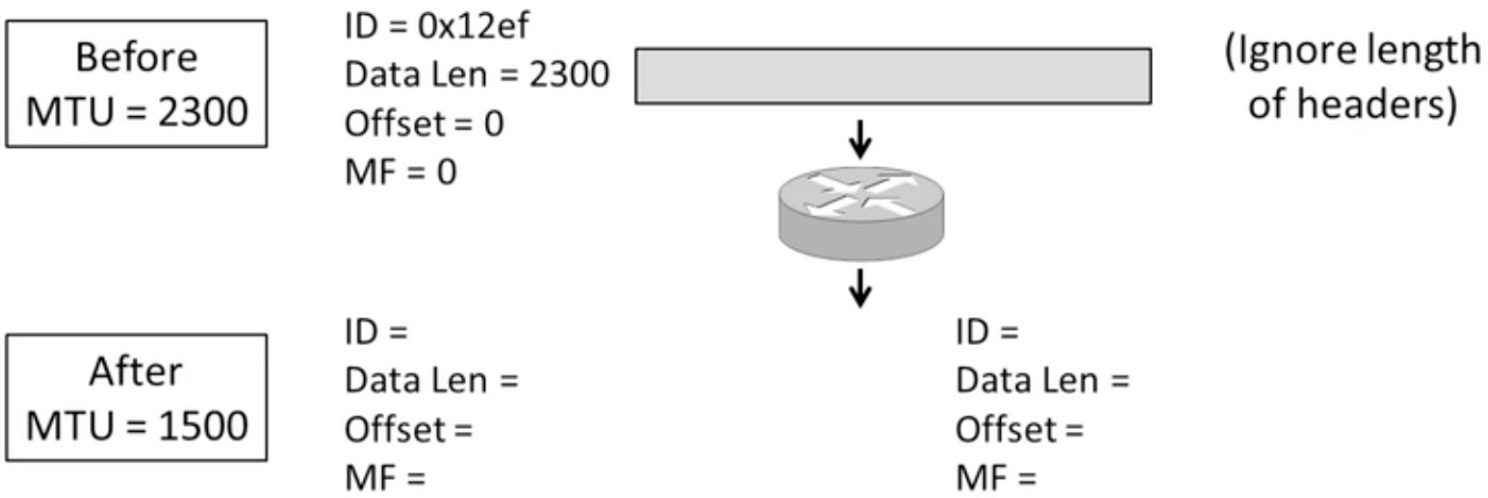

IPv4 Fragmentation (3)

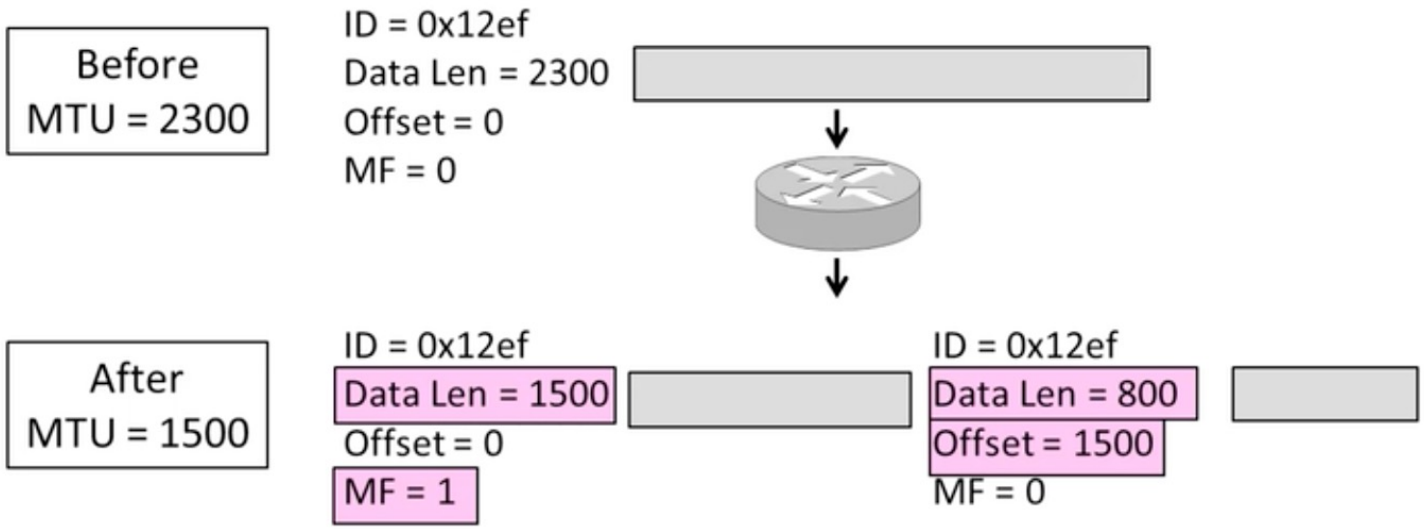

IPv4 Fragmentation (4)

- It works!
  - Allow repeated fragmentation
- But fragmentation is undesirable
  - More work for routers, hosts
  - Tends to magnify loss rate
  - Security vulnerabilities too

Path MTU Discovery

- Discover the MTU that will fit
  - So we can avoid fragmentation
  - The method in use today
- Host tests path with large packet
  - Routers provide feedback if too large; They tell host what size would have fit

Path MTU Discovery (2)

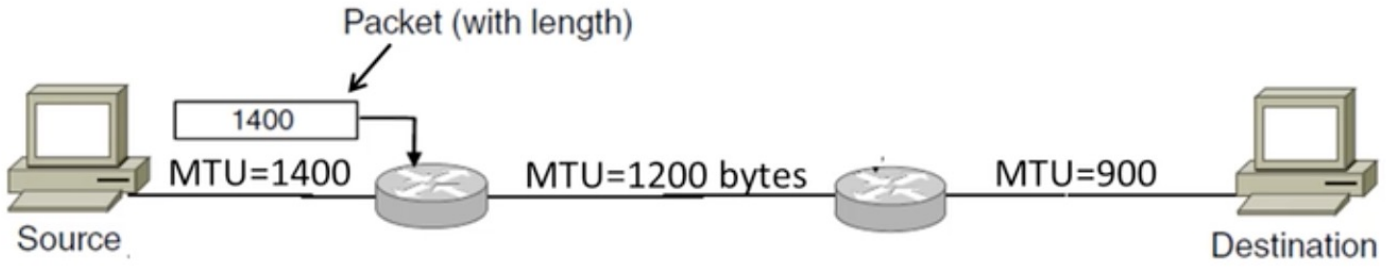

Path MTU Discovery (3)

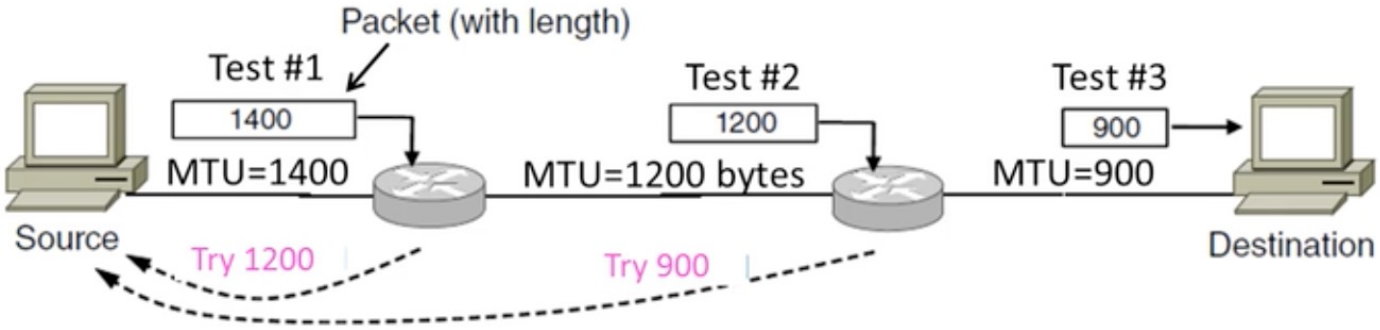

Path MTU Discovery (4)

- Process may seem involved
  - But usually quick to find right size
- Path MTU depends on the path and so can change over time
  - Search is ongoing
- Implemented with ICMP
  - Set DF (Don’t Fragment) bit in IP header to get feedback messages

Error Handling

- What happens when something goes wrong during forwarding?
  - Need to be able to find the problem

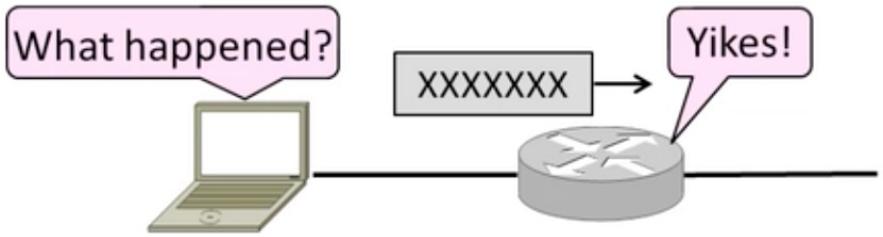

Internet Control Message Protocol

- ICMP is a companion protocol to IP
  - They are implemented together
  - Sits on top of IP (IP protocol=1)

- Provides error report and testing
  - Error is at router while forwarding
  - Also, testing that hosts can use

ICMP Errors

- When router encounters an error while forwarding:
  - It sends an ICMP error report back to the IP source address
  - It discards the problematic packet; host needs to rectify

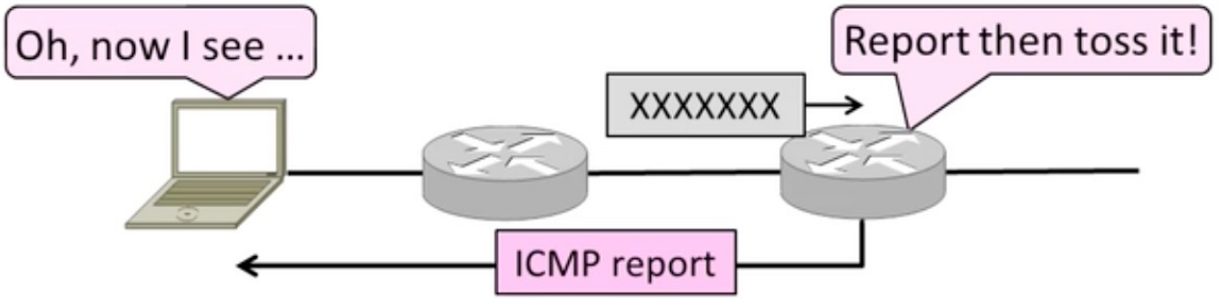

ICMP Message Format

- Each ICMP message has a type, code and checksum
- Often carry the start of the offending packet as payload
- Each message is carried in an IP packet

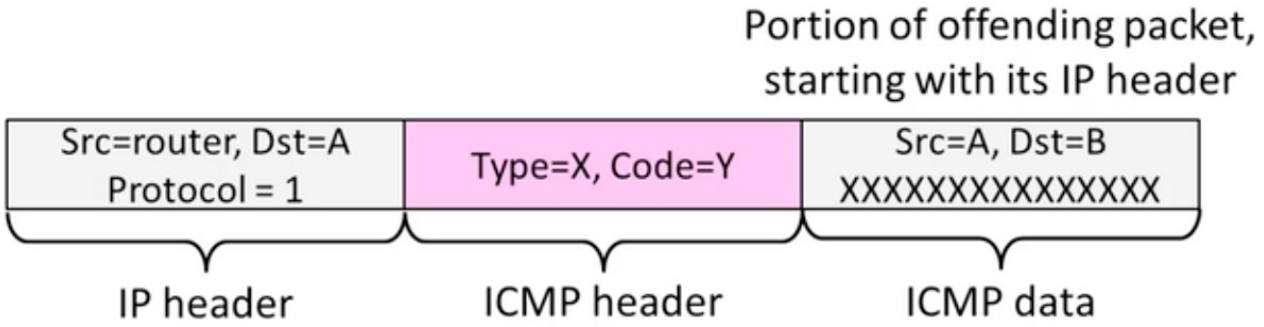

Example ICMP Message

Name	Type / Code	Usage
Dest. Unreachable (Net or Host)	3 / 0 or 1	Lack of connectivity
Dest. Unreachable (Fragment)	3 / 4	Path MTU Discovery
Time Exceeded (Transit)	11 / 0	Traceroute
Echo Request or Reply	8 or 0 / 0	Ping

Testing, not a forwarding error: Host sends Echo Request, and destination responds with an Echo Reply

Traceroute

- IP header contains TTL (Time To Live) field
  - Decremented every router hop, with ICMP error if it hits zero
  - Protects against forwarding loops

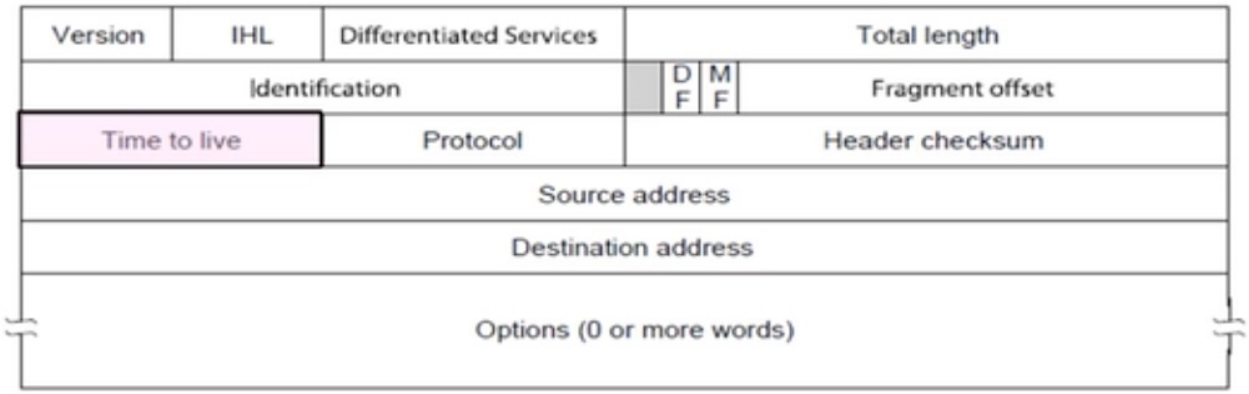

Traceroute (2)

- Traceroute repurposes TTL and ICMP functionality
  - Sends probe packets increasing TTL starting from 1
  - ICMP errors identify routers on the path

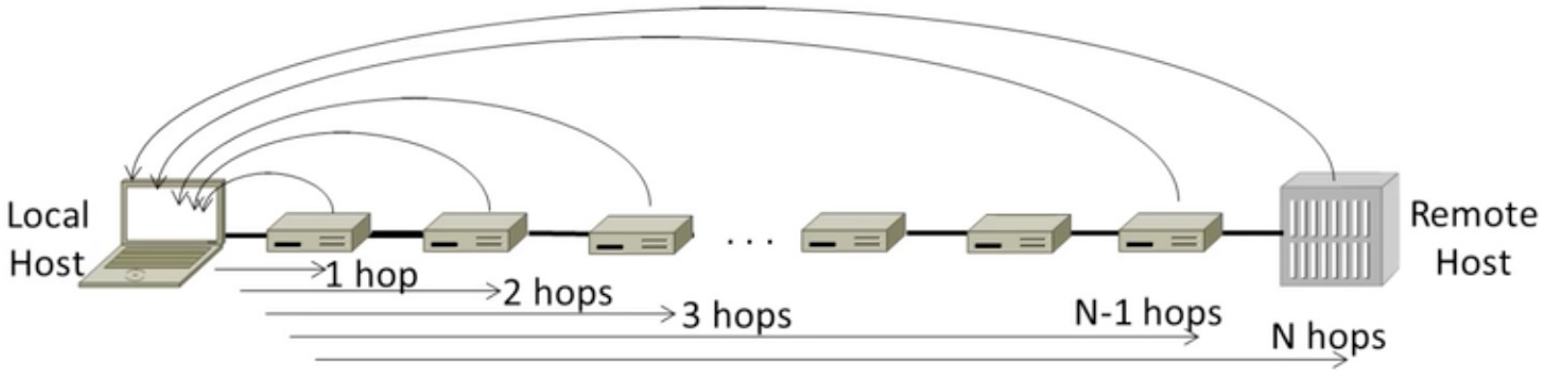

Internet Growth

- At least a billion Internet hosts and growing …
  - And we're using 32-bit addresses!
- IP version 6, the future of IPv4 that is now (still) being deployed

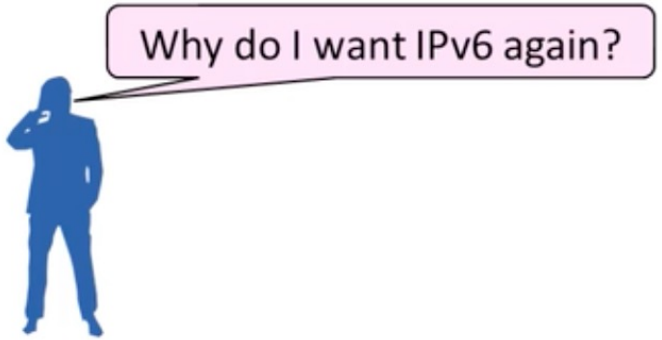

IPv6

- Feature large addresses
  - 128 bits, most of header
- New notation
  - 8 groups of 4 hex digits (16 bits)
  - Omit leading zeros, groups of zeros

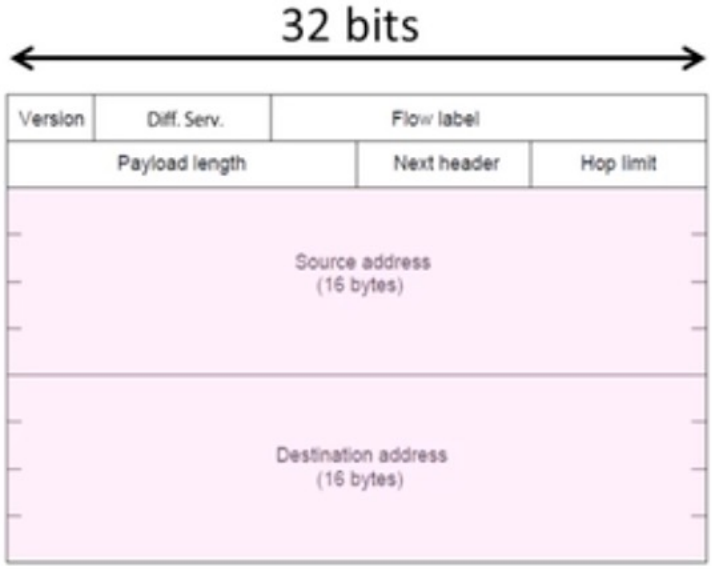

  - EX: 2001:0db8:0000:0000:0000:ff00:0042:8329
  - ->2001:db8::ff00:42:8329

IPv6 Transition

- The big problem:
  - How to deploy IPv6?
  - Fundamentally incompatible with IPv6?

- Dozens of approaches proposed
  - Dual stack (speak IPv4 and IPv6)
  - Translators (convert packets)
  - Tunnels (carry IPv6 over IPv4)

Tunneling

- Native IPv6 islands connected via IPv4
  - Tunnel carries IPv6 packets across IPv4 network

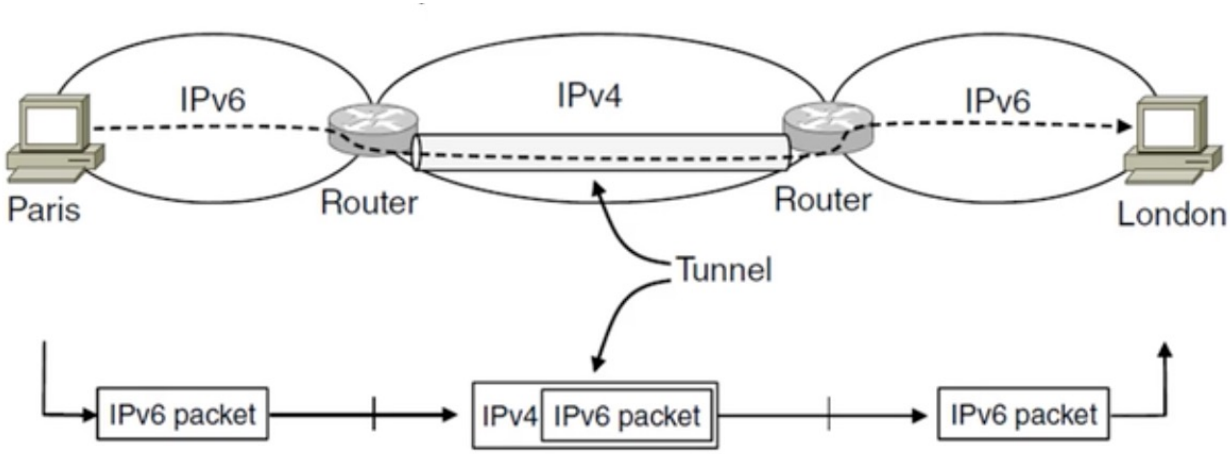

Tunneling (2)

- Tunnel acts as a single link across IPv4 network

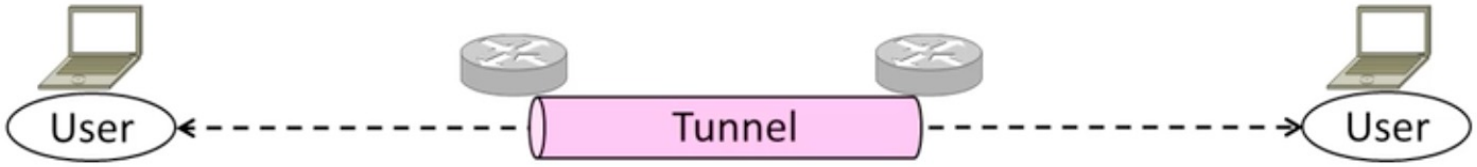

Tunneling (3)

- Tunnel acts as a single link across IPv4 network

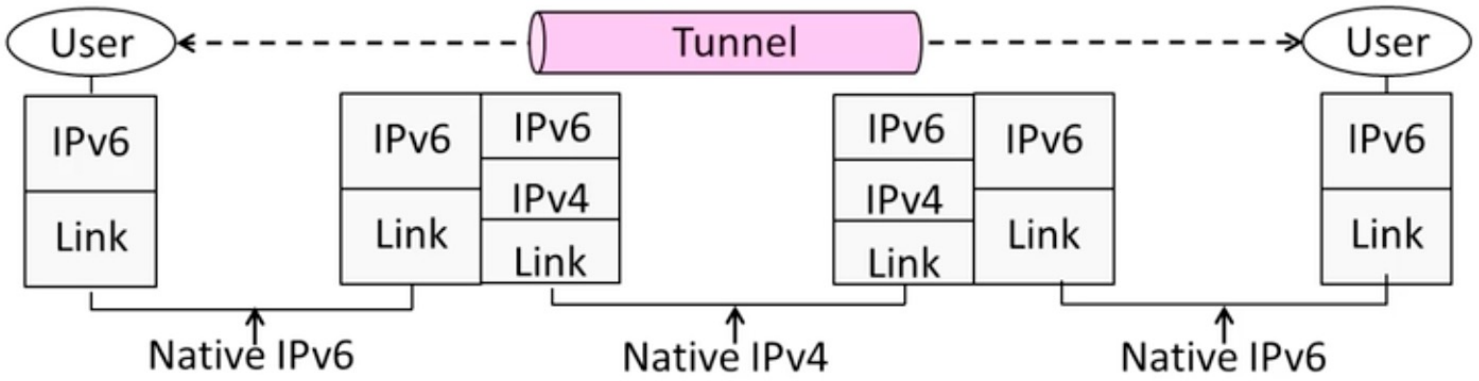

Layering Review

- Remember how layering is meant to work?
  - "Routers don't look beyond the IP header." Well …

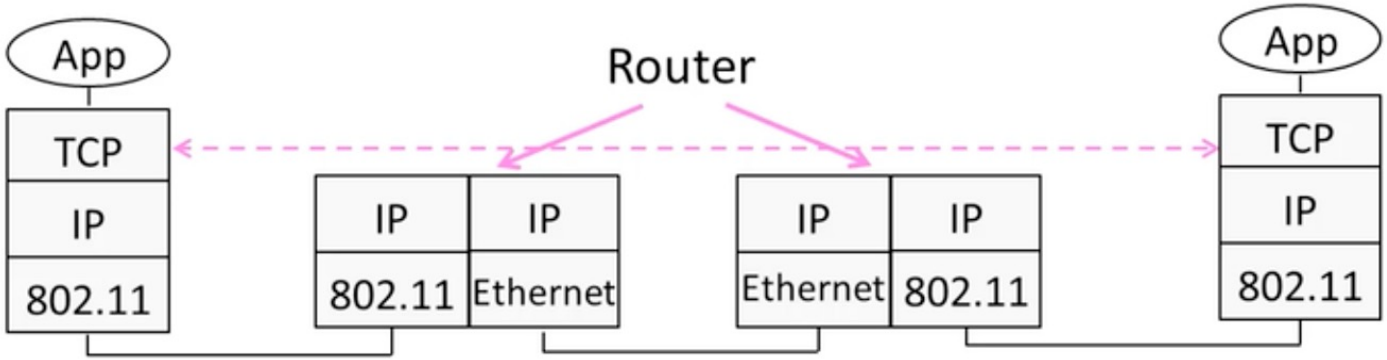

Middleboxes

- Sit “inside the network” but perform “more than IP” processing on packets to add new functionality
  - NAT box, Firewall / Intrusion Detection System

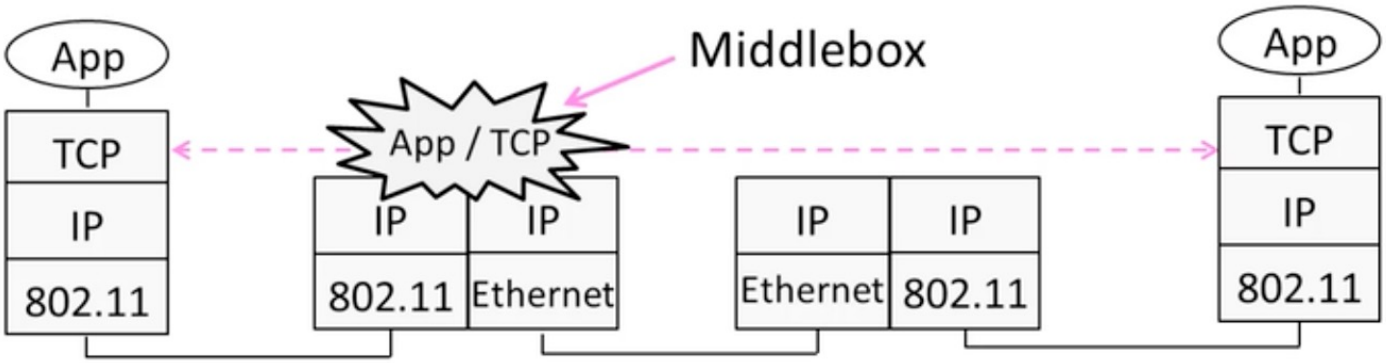

What is NAT (Network Address Translation)

- NAT box connects an internal network to an external network
  - Many internal hosts are connected using few external addresses
  - Middlebox that “translates addresses”
- Motivated by IP address scarcity
  - Controversial at first, now accepted

NAT (2)

- Common scenario:
  - Home computers use “private” IP addresses
  - NAT (in AP/firewall) connects home to ISP using a single external IP address

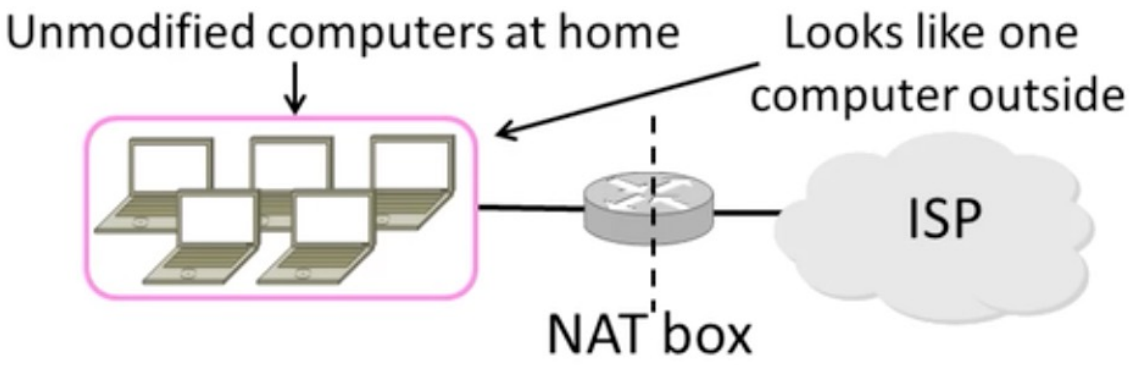

How NAT Works?

- Keeps an internal/external table
  - Typically uses IP address + TCP port
  - This is address and port translation

What host thinks	What ISP thinks
Internal IP:port	**External IP : port**
192.168.1.12 : 5523	44.25.80.3 : 1500
192.168.1.13 : 1234	44.25.80.3 : 1501
192.168.2.20 : 1234	44.25.80.3 : 1502

- Need ports to make mapping 1-1 since there are fewer external IPs

How NAT Works? (2)

- Internal -> External:
  - Look up and rewire source IP/port

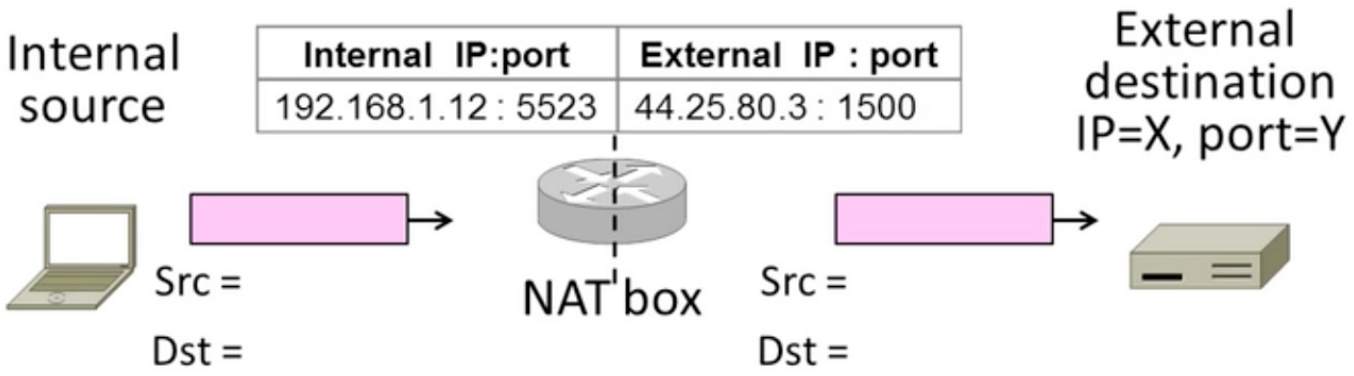

How NAT Works? (3)

- External -> Internal:
  - Look up and rewire destination IP/port

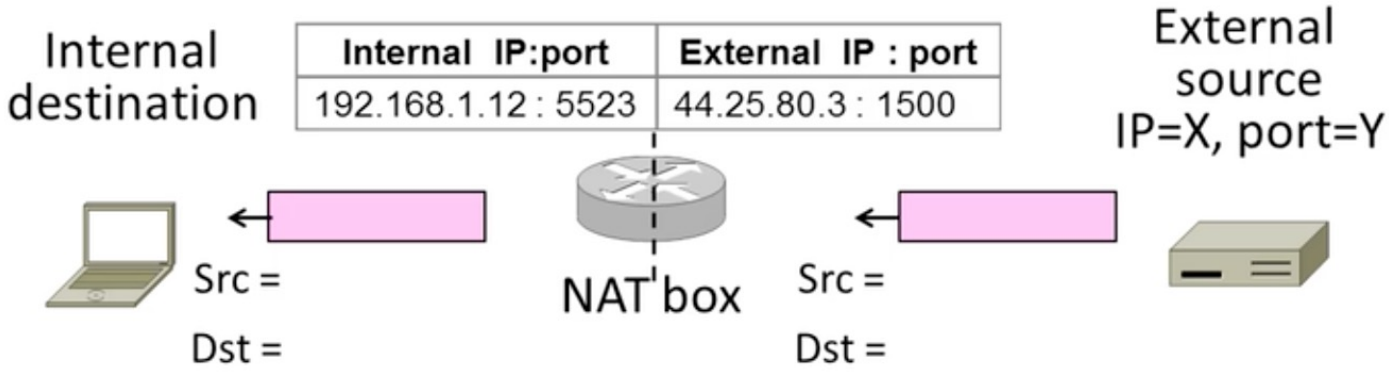

NAT Downsides

- Connectivity has been broken!
  - Can only send incoming packets after an outgoing connection is set up
  - Difficult to run servers or peer-to-peer apps (Skype) at home

- Doesn’t work so well when there are no connections (UDP apps)
- Breaks apps that unwisely expose their IP addresses (FTP)

NAT Upsides

- Relieves much IP address pressure
  - Many home hosts behind NATs
- Easy to deploy
  - Rapidly, and by you alone
- Useful functionality
  - Firewall, helps with privacy

To do

- Quiz next week
- Research reference due March 1st
- Lab1 due March 15th